


# Implicit Context Condensation for Local SWE Agents

AI Agents & Planning

**Kirill Gelvan** w\ Igor Slinko, Felix Steinbauer, Yaroslav Zharov, Gjergji Kasneci

10.12.2025




JETBRAINS

1. Motivation
2. Alternative Approaches
3. Methodology
4. Experiments
5. Discussion & Future Work
6. Conclusion


# Motivation

---


# Example




# Example




# Example



# Example




# Example



# Why bother with implicit condensation?

- Embeddings can hold lots of information
- Theoretically can be faster
- Avoids external retrieval systems



# Research Questions

RQ1: How does implicit context condensation influence the efficiency of LLM-based agents when applied to software engineering tasks?

RQ2: How does the performance of implicit context condensation on standard NLP benchmarks transfer to SWE tasks, single-shot and agentic?

# Alternative Approaches

---

- Extended context windows

- Infini-attention Transformer [1]

- ✓ 100K–1M token contexts

- ✗ New architecture + re-training

- Explicit compression

- LLM Lingua-2 (prompt pruning) [2], SlimCode (code pruning) [3]

- ✓ Model-agnostic, cheaper & faster

- ✗ Lossy by design

- Soft-tokens

- CoCoNut (soft-tokens for reasoning) [4]

- ✓ Faster reasoning via continuous vectors + inspiring

- ✗ Applied to reasoning tokens, uses RL


# The focus – Implicit compression

- ✓ Continuous, theoretically near-lossless compression
- ✓ Plug-in to existing LLMs, can be fine-tuned on coding


# Methodology

# In-Context AutoEncoder (ICAE) [5]


- Same LLM as both compressor and decoder
- Encoder (LoRA) writes full context into a few learned memory slots
- Decoder reads only memory slots + Special Token (ST = AE or LM)
- Trained first for reconstruction, then FT for tasks




# Training process & model names



# Training process & model names




# Training process & model names



# Pre-Training on general texts

- General texts dataset (SlimPajama6B[6])
- 50% AE and 50% LM task
- Train only the encoder's LoRA weights so its memory slots let the frozen LLM reconstruct or continue the text

# Fine-Tuning (incl. on trajectories)



# Experiments

---

---

---

---

# Metrics

- BLEU – reconstruction / solution quality based on n-gram precision [8]
- Resolve rate – fraction of tasks successfully solved
- Time – mean time per generated tool-call
- # Turns – interaction steps needed to reach a solution

# Pretraining

- 4x compression

## Autoencoding Reconstruction Failure Example

Original:

```
<p align="center">
<a href="https://swe-agent.com/latest/">
<strong>Documentation</strong></a>&nbsp;;
...

```

Reconstructed:

```
<p align="center">
<a href="https://swe-agent.com /agent/ latest/">
<strong>Documentation</strong></a>&nbsp;;
...

```

# Pretraining (AE scores)

Run	Checkpoint (# steps)	Compression	BLEU (mean, n=100)
Qwen3-8B (no ICAE)	18k	×1	0.867
<i>10M tokens subset</i>			
ICAE-PT	9k	×4	0.909
ICAE-PT	12k	×4	<u>0.942</u>
ICAE-PT	18k	×4	0.902
<i>1B tokens subset</i>			
ICAE-PT	9k	×4	0.936
ICAE-PT	12k (main)	×4	<b>0.964</b>
ICAE-PT	18k	×4	0.928

# SQuAD [9]

Single-step QA  
(dates, names, etc)


In meteorology, precipitation is any product of the condensation of atmospheric water vapor that falls under **gravity**. The main forms of precipitation include drizzle, rain, sleet, snow, **graupel** and hail... Precipitation forms as smaller droplets coalesce via collision with other rain drops or ice crystals **within a cloud**. Short, intense periods of rain in scattered locations are called “showers”.

What causes precipitation to fall?  
**gravity**

What is another main form of precipitation besides drizzle, rain, snow, sleet and hail?  
**graupel**


Where do water droplets collide with ice crystals to form precipitation?  
**within a cloud**

## SQuAD




# RepoQA [10]

## Single-step code reconstruction (needle in a haystack)





## RepoQA





# SWE-Bench Verified [11]

## Agentic SWE



- Observations of length >256 tokens are compressed



# Research Questions

RQ1: How does implicit context condensation **influence the efficiency** of LLM-based agents when applied to **software engineering tasks**?

# Time

Max # steps: 75


Max context len: 32,768 tokens

Name	Time (s) ↓
Naive (del long obs-s)	0.44
Baseline+FT	1.24
Baseline	1.23
ICAE-PT	1.23
ICAE-PT+FT	<b>1.12 (0.31+0.81)</b>

# # Steps

Max # steps: -


Max context len: 32,768 tokens




# Research Questions

RQ2: How does the **performance** of implicit context condensation on standard **NLP** benchmarks **transfer to SWE** tasks, **single-shot** and **agentic**?

# BLEU



# BLEU vs Resolve Rate



# Discussion & Future Work

---

# Reconstruction fidelity & cascading errors

- Slightly imperfect BLEU → small bugs can accumulate in long trajectories

Future steps:

- Improve pretraining BLEU & tasks
- Explore other compression ratios

## Autoencoding Reconstruction Failure Example

Original:

```
<p align="center">
<a href="https://swe-agent.com/latest/">
<strong>Documentation</strong></a>&nbsp;;
...
```

Reconstructed:


```
<p align="center">
<a href="https://swe-agent.com /agent/ latest/">
<strong>Documentation</strong></a>&nbsp;;
...
```

# LoRA does not work with trajectories-FT

- Low-rank approximations may learn teacher **style**, not **robust policies**

Future steps:

- Use full-model training
- Explore RL-style optimization with direct rewards from SWE-bench




# Information preserving for multi-turn interactions

- Per-step compression + **single-step loss** → throw away details that matter **many steps later**

Future steps:

- Train over multiple steps
- Use trajectory-level objectives (e.g., RL) to retain information critical for later steps




# Model size and reasoning

- Only 8B-sized model **without reasoning** might be insufficient for multi-turn agentic SWE

Future steps:

- Train model of sizes 32B+
- Try FT on teacher's model reasoning as well



# Conclusion

---

- RQ1 – Efficiency: ICAE-based compression **enables longer trajectories** (by  $\approx 40\%$ ) and **tool-call speedups** (by  $\approx 10\%$ ) under a fixed context window.
- RQ2 – Transferability: Compression **transfers positively** to **QA** and **single-shot SWE tasks**, but **degrades agentic SWE performance**.

Let's discuss! 🤔

# References

- [1] T. Munkhdalai, M. Faruqui, and S. Gopal, "Leave no context behind: Efficient infinite context transformers with infini-attention," arXiv:2404.07143, 2024.
- [2] Z. Pan et al., "LLMLingua-2: Data distillation for efficient and faithful task-agnostic prompt compression," arXiv:2403.12968, 2024.
- [3] Y. Wang et al., "Natural is the best: Model-agnostic code simplification for pre-trained large language models," Proc. ACM on Software Engineering 1(FSE), pp. 586–608, 2024.
- [4] CoCoNut – not listed in the thesis bibliography; please insert your preferred CoCoNut citation here.
- [5] T. Ge et al., "In-context autoencoder for context compression in a large language model," arXiv:2307.06945, 2023.
- [6] M. Weber et al., "RedPajama: An open dataset for training large language models" (used as the SlimPajama-6B pretraining corpus), Advances in Neural Information Processing Systems 37, 2024.
- [7] J. Yang et al., "SWE-Smith: Scaling data for software engineering agents," arXiv:2504.21798, 2025.
- [8] K. Papineni et al., "BLEU: A method for automatic evaluation of machine translation," in Proc. 40th Annual Meeting of the ACL, pp. 311–318, 2002.
- [9] P. Rajpurkar et al., "SQuAD: 100,000+ questions for machine comprehension of text," arXiv:1606.05250, 2016.
- [10] J. Liu et al., "RepoQA: Evaluating long context code understanding," arXiv:2406.06025, 2024.
- [11] C. E. Jimenez et al., "SWE-bench: Can language models resolve real-world GitHub issues?" arXiv:2310.06770, 2023.

# Backup Slides

Name in our paper	Encoder	Decoder	Resolved (/500) ↑	Time (s) ↓
<i>Naive Baselines</i>				
—	del long obs-s	Qwen	1	0.44
—	del all obs-s	Qwen	0	0.39
<i>Baselines</i>				
—	—	Qwen (Full-FT)	<b>86</b>	1.24
Baseline+FT	—	Qwen (LoRA-FT)	10	1.24
Baseline	—	Qwen	<u>19.4 ± 6.5</u>	1.23
<i>ICAE Compression (ours)</i>				
ICAE-PT	ICAE (LoRA-PT)	Qwen	2	1.23
ICAE-PT+FT	ICAE (LoRA-PT & LoRA-FT)	Qwen	$7.8 \pm 2.59$	<b>1.12 (0.31+0.81)</b>
—	ICAE (LoRA-PT & LoRA-FT)	Qwen (LoRA-FT)	10	<u>1.13</u>

---- BEGIN FUNCTION #1: bash ----

Description: Execute a bash command in the terminal.

Parameters:

- (1) command (string, required): The bash command to execute. Can be empty to view additional logs when previous exit code is '-1'. Can be `ctrl+c` to interrupt the currently running process.

---- END FUNCTION #1 ----

---- BEGIN FUNCTION #2: submit ----

Description: Finish the interaction when the task is complete OR if the assistant cannot proceed further with the task.

No parameters are required for this function.

---- END FUNCTION #2 ----

---- BEGIN FUNCTION #3: str\_replace\_editor ----

Description: Custom editing tool for viewing, creating and editing files

- \* State is persistent across command calls and discussions with the user
- \* If `path` is a file, `view` displays the result of applying `cat -n`. If `path` is a directory, `view` lists non-hidden files and directories up to 2 levels deep
- \* The `create` command cannot be used if the specified `path` already exists as a file
- \* If a `command` generates a long output, it will be truncated and marked with ``
- \* The `undo\_edit` command will revert the last edit made to the file at `path`

Notes for using the `str\_replace` command:

- \* The `old\_str` parameter should match EXACTLY one or more consecutive lines from the original file. Be mindful of whitespaces!
- \* If the `old\_str` parameter is not unique in the file, the replacement will not be performed. Make sure to include enough context in `old\_str` to make it unique
- \* The `new\_str` parameter should contain the edited lines that should replace the `old\_str`

Parameters:

- (1) command (string, required): The commands to run. Allowed options are: `view`, `create`, `str\_replace`, `insert`, `undo\_edit`.

Allowed values: [`view`, `create`, `str\_replace`, `insert`, `undo\_edit`]

- (2) path (string, required): Absolute path to file or directory, e.g., `/repo/file.py` or `/repo`.

- (3) file\_text (string, optional): Required parameter of `create` command, with the content of the file to be created.

- (4) old\_str (string, optional): Required parameter of `str\_replace` command containing the string in `path` to replace.

- (5) new\_str (string, optional): Optional parameter of `str\_replace` command containing the new string (if not given, no string will be added). Required parameter of `insert` command containing the string to insert.

- (6) insert\_line (integer, optional): Required parameter of `insert` command. The `new\_str` will be inserted AFTER the line `insert\_line` of `path`.

- (7) view\_range (array, optional): Optional parameter of `view` command when `path` points to a file. If none is given, the full file is shown. If provided, the file will be shown in the indicated line number range, e.g., [11, 12] will show lines 11 and 12. Indexing at 1 to start. Setting [start\_line, -1] shows all lines from start\_line to the end of the file.

---- END FUNCTION #3 ----